

CNN Basics

Shubhra Aich
s.aich.72@gmail.com

January 17, 2018

Outline

RGB image basics

ANN vs CNN

Operations in CNN

- Convolution

- Pooling



- Rectification

- Dropout


RGB image basics

RGB image consists of 3 channels or planes, namely Red, Green and Blue.

We see the combined effect of these channels.



Artificial vs Convolutional NN



Operations in CNN

- ▶ Convolution
- ▶ Rectification
- ▶ Pooling
- ▶ Dropout

Convolution ...(1)

Input Volume (+pad 1) (7x7x3)

$x[:, :, 0]$						
0	0	0	0	0	0	0
0	0	2	2	0	0	0
0	1	0	2	0	1	0
0	1	2	1	1	2	0
0	0	1	1	1	1	0
0	2	2	1	0	2	0
0	0	0	0	0	0	0

$x[:, :, 1]$						
0	0	0	0	0	0	0
0	2	2	0	1	1	0
0	0	1	2	1	0	0
0	0	0	0	0	2	0
0	0	2	0	1	2	0
0	0	1	2	0	0	0
0	0	0	0	0	0	0

$x[:, :, 2]$						
0	0	0	0	0	0	0
0	2	2	0	2	2	0
0	2	1	1	1	2	0
0	2	0	1	0	0	0
0	0	1	2	0	2	0
0	1	2	2	2	1	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

$w0[:, :, 0]$		
-1	-1	-1
-1	1	1
1	1	-1

$w0[:, :, 1]$		
-1	-1	0
1	-1	1
1	1	0

$w0[:, :, 2]$		
1	-1	1
0	1	0
-1	-1	-1

Bias $b0$ (1x1x1)

$b0[:, :, 0]$		
1		

Filter W1 (3x3x3)

$w1[:, :, 0]$		
0	1	0
0	1	-1
-1	1	0

$w1[:, :, 1]$		
0	0	-1
-1	0	-1
1	-1	0

$w1[:, :, 2]$		
1	1	0
1	-1	-1
1	-1	0

Bias $b1$ (1x1x1)

$b1[:, :, 0]$		
0		

Output Volume (3x3x2)

$o[:, :, 0]$		
3	6	2
2	-2	0
7	-5	-3

$o[:, :, 1]$		
-9	0	0
-1	3	3
-6	1	6

Convolution ... (2)

Input Volume (+pad 1) (7x7x3)

$x[:, :, 0]$

0	0	0	0	0	0	0
0	0	2	2	0	0	0
0	1	0	2	0	1	0
0	1	2	1	1	2	0
0	0	1	1	1	1	0
0	2	2	1	0	2	0
0	0	0	0	0	0	0

$x[:, :, 1]$

0	0	0	0	0	0	0
0	2	2	0	1	1	0
0	0	1	2	1	0	0
0	0	0	0	0	2	0
0	0	2	0	1	2	0
0	0	1	2	0	0	0
0	0	0	0	0	0	0

$x[:, :, 2]$

0	0	0	0	0	0	0
0	2	2	0	2	2	0
0	2	1	1	1	2	0
0	2	0	1	0	0	0
0	0	1	2	0	2	0
0	1	2	2	2	1	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

$w0[:, :, 0]$

-1	-1	-1
-1	1	1
1	1	-1

$w0[:, :, 1]$

-1	-1	0
1	-1	1
1	1	0

$w0[:, :, 2]$

1	-1	1
0	1	0
-1	-1	-1

Bias b0 (1x1x1)

$b0[:, :, 0]$

1

Filter W1 (3x3x3)

$w1[:, :, 0]$

0	1	0
0	1	-1
-1	1	0

$w1[:, :, 1]$

0	0	-1
-1	0	-1
1	-1	0

$w1[:, :, 2]$

1	1	0
1	-1	-1
1	-1	0

Bias b1 (1x1x1)

$b1[:, :, 0]$

0

Output Volume (3x3x2)

$o[:, :, 0]$

3	6	2
2	-2	0
7	-5	-3

$o[:, :, 1]$

-9	0	0
-1	3	3
-6	1	6

Convolution ... (3)

Input Volume (+pad 1) (7x7x3)

$x[:, :, 0]$

0	0	0	0	0	0	0
0	0	2	2	0	0	0
0	1	0	2	0	1	0
0	1	2	1	1	2	0
0	0	1	1	1	1	0
0	2	2	1	0	2	0
0	0	0	0	0	0	0

$x[:, :, 1]$

0	0	0	0	0	0	0
0	2	2	0	1	1	0
0	0	1	2	1	0	0
0	0	0	0	0	2	0
0	0	2	0	1	2	0
0	0	1	2	0	0	0
0	0	0	0	0	0	0

$x[:, :, 2]$

0	0	0	0	0	0	0
0	2	2	0	2	2	0
0	2	1	1	1	2	0
0	2	0	1	0	0	0
0	0	1	2	0	2	0
0	1	2	2	2	1	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

$w0[:, :, 0]$

-1	-1	-1
-1	1	1
1	1	-1

$w0[:, :, 1]$

-1	-1	0
1	-1	1
1	1	0

$w0[:, :, 2]$

1	1	1
0	1	0
-1	-1	-1

Bias b0 (1x1x1)

$b0[:, :, 0]$

1

Filter W1 (3x3x3)

$w1[:, :, 0]$

0	1	0
0	1	-1
-1	1	0

$w1[:, :, 1]$

0	0	-1
-1	0	-1
1	-1	0

$w1[:, :, 2]$

1	1	0
1	-1	-1
1	-1	0

Bias b1 (1x1x1)

$b1[:, :, 0]$

0

Output Volume (3x3x2)

$o[:, :, 0]$

3	6	2
2	-2	0
7	-5	-3

$o[:, :, 1]$

-9	0	0
-1	3	3
-6	1	6

Convolution ... (4)

Input Volume (+pad 1) (7x7x3)

$x[:, :, 0]$

0	0	0	0	0	0	0
0	0	2	2	0	0	0
0	1	0	2	0	1	0
0	1	2	1	1	2	0
0	0	1	1	1	1	0
0	2	2	1	0	2	0
0	0	0	0	0	0	0

$x[:, :, 1]$

0	0	0	0	0	0	0
0	2	2	0	1	1	0
0	0	1	2	1	0	0
0	0	0	0	0	2	0
0	0	2	0	1	2	0
0	0	1	2	0	0	0
0	0	0	0	0	0	0

$x[:, :, 2]$

0	0	0	0	0	0	0
0	2	2	0	2	2	0
0	2	1	1	1	2	0
0	2	0	1	0	0	0
0	0	1	2	0	2	0
0	1	2	2	2	1	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

$w0[:, :, 0]$

-1	-1	-1
-1	1	1
1	1	-1

$w0[:, :, 1]$

-1	-1	0
1	-1	1
1	1	0

$w0[:, :, 2]$

1	-1	1
0	1	0
-1	-1	-1

Bias b0 (1x1x1)

$b0[:, :, 0]$

1

Filter W1 (3x3x3)

$w1[:, :, 0]$

0	1	0
0	1	-1
-1	1	0

$w1[:, :, 1]$

0	0	-1
-1	0	-1
1	-1	0

$w1[:, :, 2]$

1	1	0
1	-1	-1
1	-1	0

Bias b1 (1x1x1)

$b1[:, :, 0]$

0

Output Volume (3x3x2)

$o[:, :, 0]$

3	6	2
2	-2	0
7	-5	-3

$o[:, :, 1]$

-9	0	0
-1	3	3
-6	1	6

Convolution ... (5)

Input Volume (+pad 1) (7x7x3)

$x[:, :, 0]$

0	0	0	0	0	0	0
0	0	2	2	0	0	0
0	1	0	2	0	1	0
0	1	2	1	1	2	0
0	0	1	1	1	1	0
0	2	2	1	0	2	0
0	0	0	0	0	0	0

$x[:, :, 1]$

0	0	0	0	0	0	0
0	2	2	0	1	1	0
0	0	1	2	1	0	0
0	0	0	0	0	2	0
0	0	2	0	1	2	0
0	0	1	2	0	0	0
0	0	0	0	0	0	0

$x[:, :, 2]$

0	0	0	0	0	0	0
0	2	2	0	2	2	0
0	2	1	1	1	2	0
0	2	0	1	0	0	0
0	0	1	2	0	2	0
0	1	2	2	2	1	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

$w0[:, :, 0]$

-1	1	-1
-1	1	1
1	1	-1

$w0[:, :, 1]$

-1	-1	0
1	-1	1
1	1	0

$w0[:, :, 2]$

1	-1	1
0	1	0
-1	-1	-1

Bias b0 (1x1x1)

$b0[:, :, 0]$

1

Filter W1 (3x3x3)

$w1[:, :, 0]$

0	1	0
0	1	-1
-1	1	0

$w1[:, :, 1]$

0	0	-1
-1	0	-1
1	-1	0

$w1[:, :, 2]$

1	1	0
1	-1	-1
1	-1	0

Bias b1 (1x1x1)

$b1[:, :, 0]$

0

Output Volume (3x3x2)

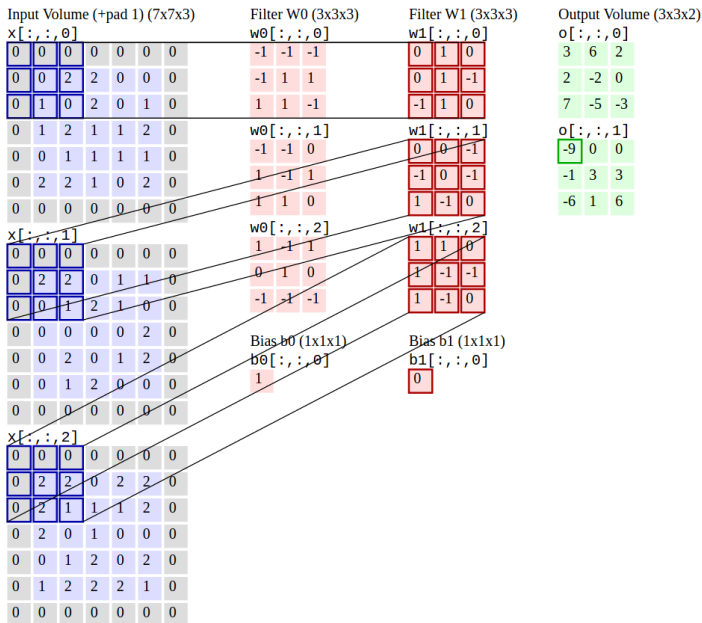
$o[:, :, 0]$

3	6	2
2	-2	0
7	-5	-3

$o[:, :, 1]$

-9	0	0
-1	3	3
-6	1	6

Convolution ... (6)



Equations of Convolution

w_i = weights of the filter


x_i = outputs of the previous layer

$$z = \sum_i w_i x_i + b = w.x + b \equiv w * x + b$$

$$a = f(z)$$


f = activation / rectification / squashing function

Pooling




- ▶ Max-pooling.
- ▶ Average pooling.

Pooling



- ▶ Max-pooling.
- ▶ Average pooling.
- ▶ Min-pooling

Pooling



- ▶ Max-pooling.
- ▶ Average pooling.
- ▶ Min-pooling ☠

Why do we use Convolution and Pooling? ...(1)

CNN combines 3 architectural ideas to deal with shift, scale and distortion.

- ▶ Local receptive fields/filters
 - ▶ The idea of local filters dates back to Hubel-Wiesel's discovery of locally sensitive, orientation selective neurons in the cat's visual cortex.
 - ▶ It helps to extract elementary visual features, such as edges, corners, end-points etc.
- ▶ Shared weights
 - ▶ Distortions or shifts of the input causes the position of salient features to vary.
 - ▶ Elementary feature detectors that are useful on one region of the image are likely to be useful across the entire image.
 - ▶ This knowledge can be applied by forcing a set of units, whose receptive fields are located at different places on the to be identical.
 - ▶ These features are combined by subsequent layers to detect more meaningful features.

Why do we use Convolution and Pooling? ...(2)

► Sub-sampling/Pooling

- Once a feature is detected, its exact location becomes less important, only approximate position relative to other features is relevant.
- This precise position is even harmful they vary from image to image.
- The simplest way to reduce this precision is sub-sampling or pooling.


Rectification

- ▶ Linear neuron, $y = wx + b$
- ▶ Rectified linear neuron (ReLU), $y = \begin{cases} wx + b, & \text{if } x > 0 \\ 0, & \text{otherwise} \end{cases}$
- ▶ Sigmoid/logistic neuron, $y = \frac{1}{1 + \exp(-\alpha(wx + b))}$

Choice of Rectification Functions

- ▶ Backpropagation is the most popular learning algorithm for a network comprising real-valued units.
- ▶ This learning algorithm adjusts the weights in each iteration in proportion to the derivative of the final error function w.r.t. the weights.
- ▶ To facilitate learning, rectification functions are expected to have more or less smooth derivatives.
- ▶ Linear neurons can map only linear input-output relationships.
- ▶ Linear neurons are unbounded both below and above.
- ▶ Sigmoid activation works well for simple ANN. However, in deeper models like CNN, tiny error-gradients of sigmoid units make learning horrendously slow.

Dropout



- ▶ In each iteration of learning, we randomly omit each hidden unit with a prespecified probability.
- ▶ Using dropout, for a network comprising single hidden layer with H hidden units, in each iteration, we randomly sample from 2^H different architectures.
- ▶ All architectures share weights.
- ▶ Dropout prevents overfitting.
- ▶ In test time, the full network is used with parameter values halved.